

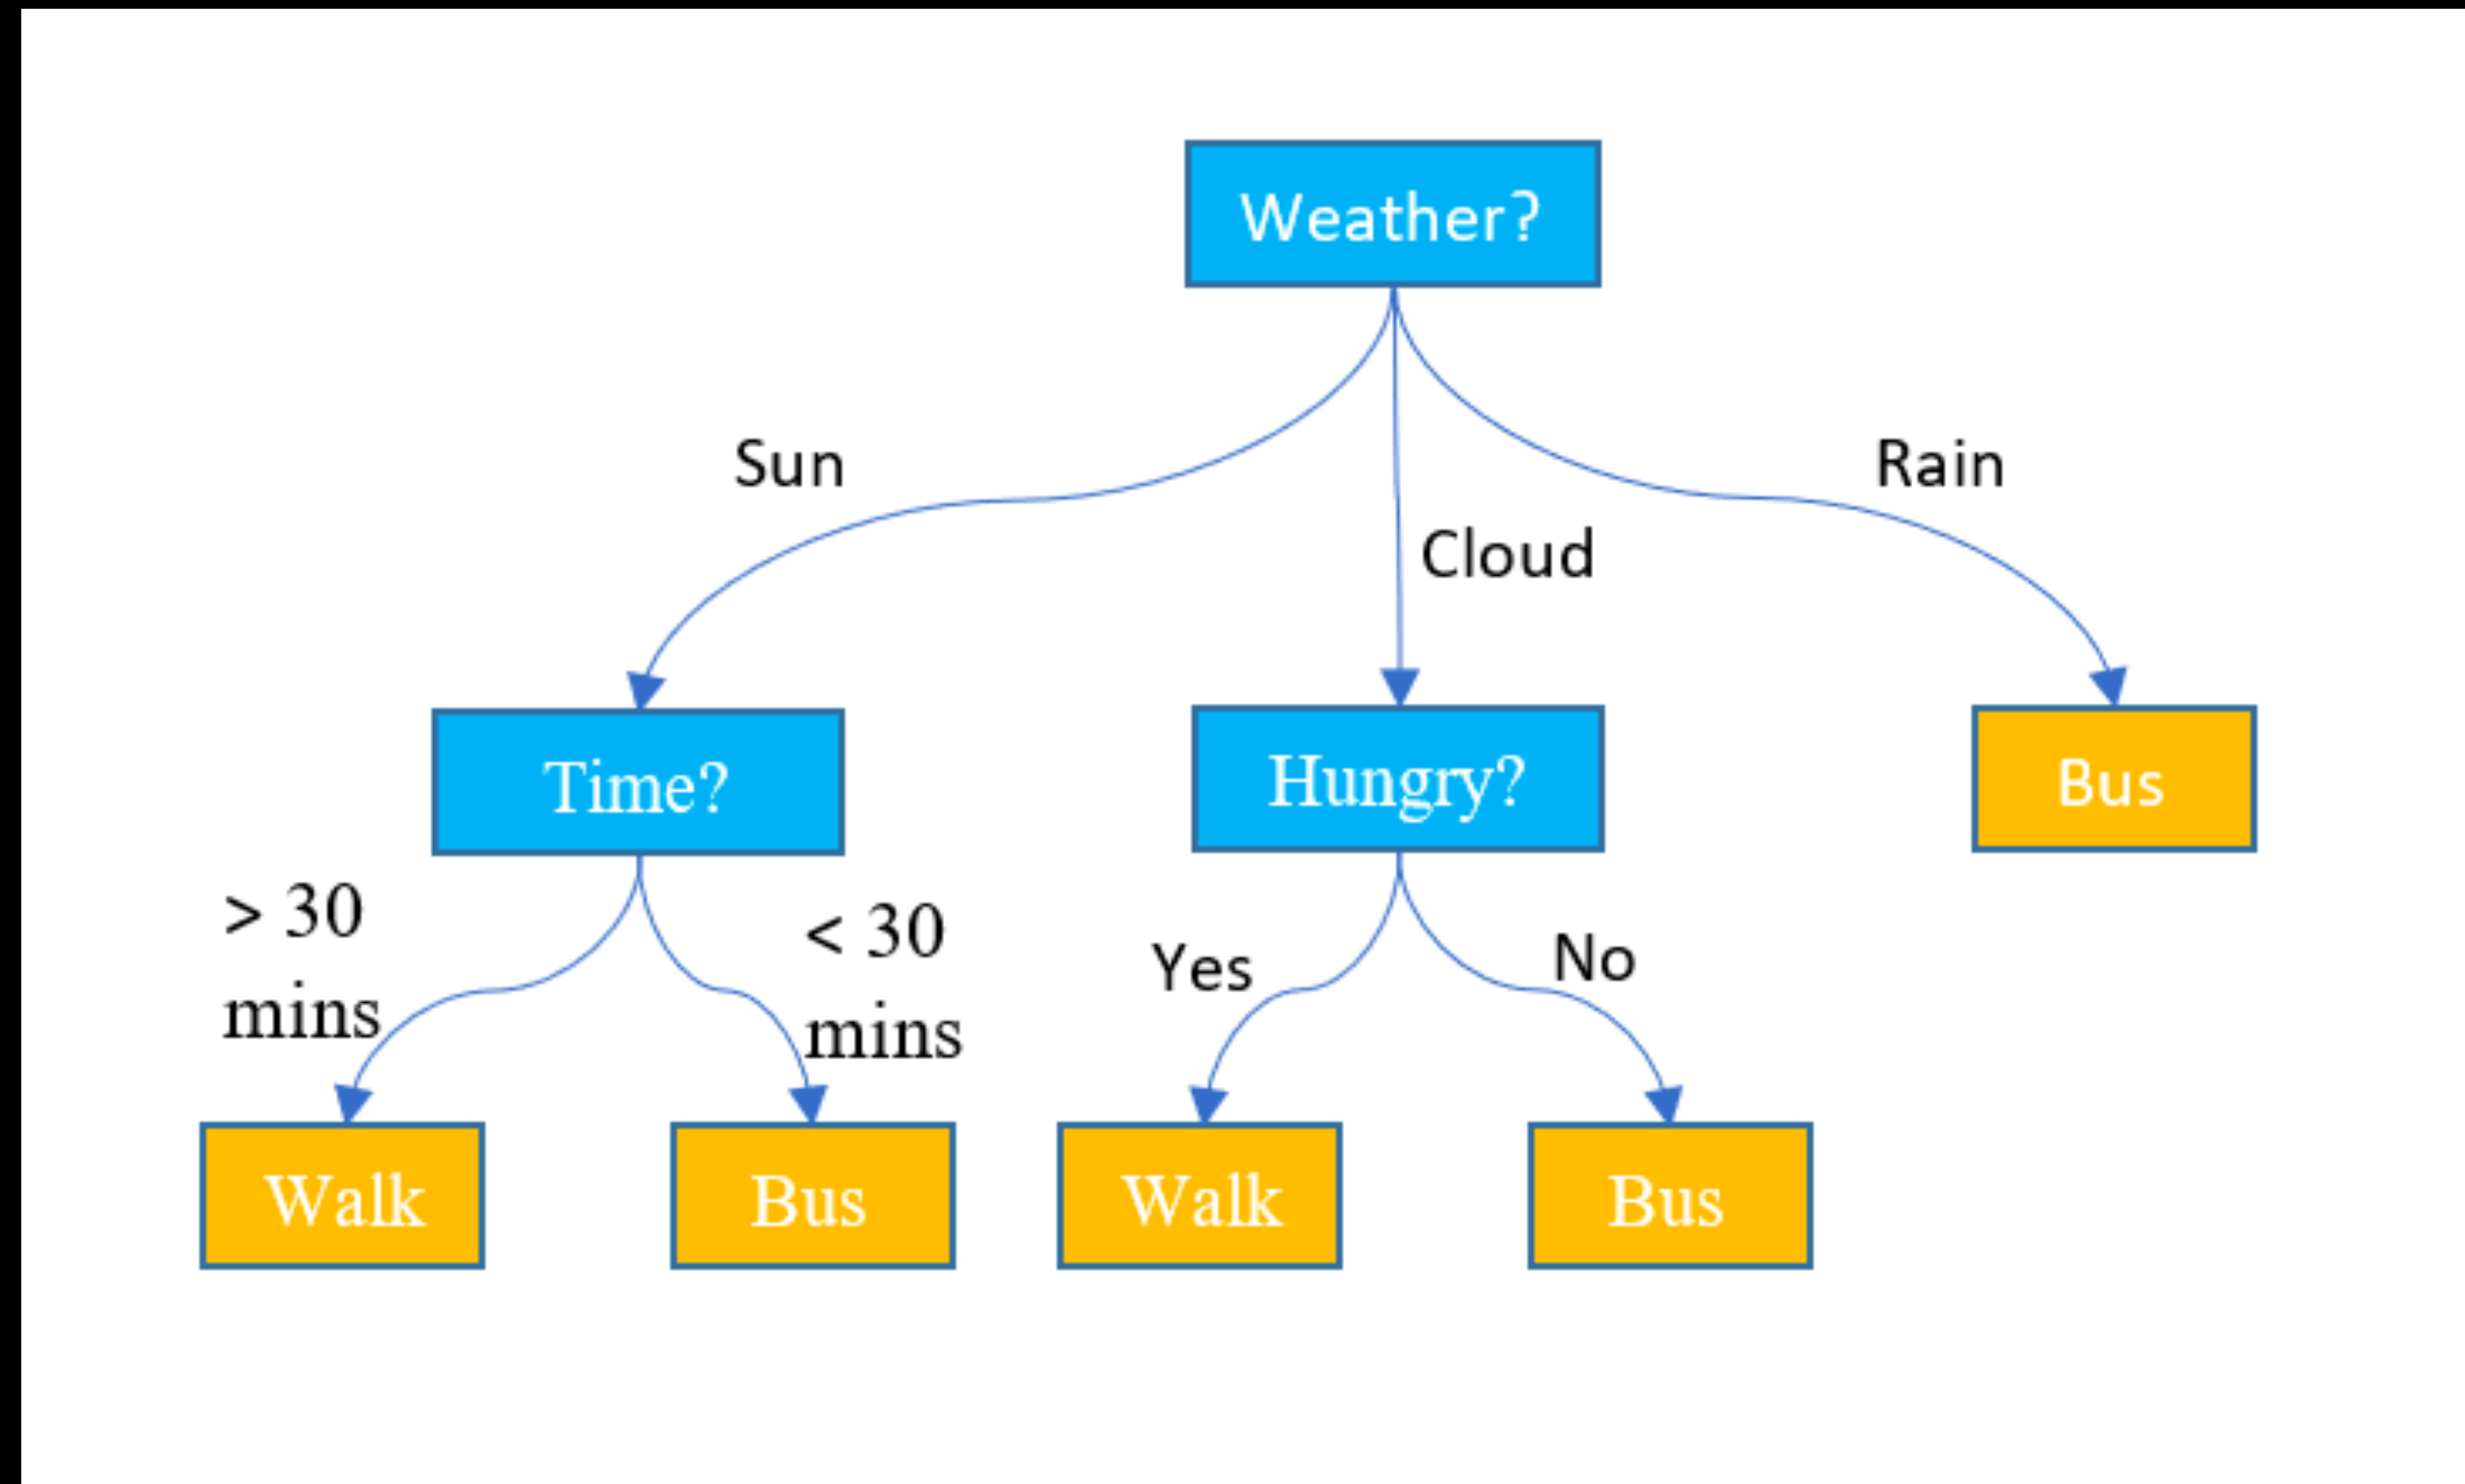
# Recommendation Systems

**Ranking** algorithms

Eugeny Malyutin / Sergey Dudorov

# Ranking algorithms

## Tree reminder



# Ranking algorithms

## Splitting criteria:

- Split:  $[x_j < t]$

- Loss criteria:  $Q(X_m, j, t) = \frac{|X_l|}{|X_M|} H(X_l) + \frac{|X_r|}{|X_m|} H(X_m)$

- Information criteria:


- Regression:  $H(x) = 1/|X_m| \sum (y_i - \bar{y}(X))^2$

- Gini:  $p_k = \frac{1}{|X|} \sum_{i \in X} [y_i = k]$  and  $H(x) = \sum_{k=1} p_k(1 - p_k)$

- Entropy:  $H(X) = \sum_{k=1}^K p_k \ln(p_k)$

# Ranking algorithms


What's went wrong?



# Ranking algorithms


What's went wrong?

- Stop when:
  - Max length
  - Min elements
  - Pruning




# Ranking algorithms

What's went wrong?



Overfitted



Overfitted and unstable

# Ranking algorithms

## Error decomposition

- What is?



# Ranking algorithms

## Error decomposition

- **Error:**
  - **Noise:** world's imperfectness measure, exists even for ideal model on ideal data
  - **Bias:** deviation of concrete model from ideal one (averaged over all data sets)
  - **Variance:** dispersion of models answers caused by different datasets




Fig. 1: Graphical Illustration of bias-variance trade-off , Source: Scott Fortmann-Roe., Understanding Bias-Variance Trade-off



# Ranking algorithms

## Composition idea

- Takes N different trees
- **Averages** answers:
  - **Regression:**  $a(x) = \sum_{1}^N b_n(x)$
  - **Classification:** most popular answer
- **PROBLEMS?**

# Ranking algorithms

## Error decomposition

- **Trees:**
  - High variance
  - Low bias
- **Linear algorithms:**
  - Low variance
  - High bias
- **Composition:**
  - Same bias
  - $\text{Variance} = 1/N (\text{algorithm\_variance}) + \text{correlation}$
  - Decrease variance **N times** in case of **independent algorithms** — that's why **we need randomisation**

# Ranking algorithms

## Randomisation

- **Bagging (bootstrap):** learn algorithms on random subsamples from train set. Less subsamples — more randomised trees.
- **Random sub-spaces:** random sub-set of features for each tree
- **Extreme randomised trees:** choose random subset of features on each split

# Ranking algorithms

## Summary

- **Cons:**
  - Powerful
  - Easy to parallel
  - With bootstrap 1/3 of data samples (for concrete tree) wasn't in train set — we can estimate test metrics on them (Out-of-bag score)
- **Cons:**
  - More trees -> more computational resources
  - Undirected search

# Ranking algorithms

## Boosting idea

- $b_0(x)$  — initial algorithm (zero, mean class, mean value)

- $a_m(x) = \sum_{i=1}^m b_i(x)$  — step of composition

- $F = \sum_{i=1}^N L(y_i, a_{m-1}(x_i) + b_i(x_i)) \rightarrow \min_b$  — optimisation

- $s = (s_1, s_2 \dots s_N)$  — displacements vector then

$$F = \sum_{i=1}^N L(y_i, a_{m-1}(x_i) + s_i) \rightarrow \min_s$$

- Optimal shift —  $\nabla F = \left[ \frac{dF}{da_{m-1}}(x_i) \right]_{i=1}^N = \left[ \frac{\sum_{i=1}^N L(y_i, a_{m-1}(x_i))}{da_{m-1}} \right]_{i=1}^N = \left[ \frac{dL(y_i, a_{m-1})}{da_{m-1}} x_i \right]_{i=1}^N$

# Ranking algorithms

## Boosting idea

- $a_m(x) = \sum_{i=1}^m b_i(x)$  – step of composition,  $F = \sum_{i=1}^N L(y_i, a_{m-1}(x_i) + b_i(x_i)) \rightarrow \min_b$  – optimisation
- Optimal shift
$$s_i = - \nabla F = \left[ \frac{dF}{da_{m-1}}(x_i) \right]_{i=1}^N = \left[ \frac{\sum_{i=1}^N L(y_i, a_{m-1}(x_i))}{da_{m-1}}(x_i) \right]_{i=1}^N = \left[ \frac{dL(y_i, a_{m-1})}{da_{m-1}}(x_i) \right]_{i=1}^N$$
- But wait,  $s_i$  is not an algorithm, it's a vector of numbers!
- Yep, but we can learn out algorithm  $b_m(x_i) \rightarrow s_i$



# Ranking algorithms

## Summary

- Powerful (extremely powerful)
- Allows all this tricks of gradient algorithms: learning rate, decay, ...
- Hard to interpret (~)
- Works mostly with weak algorithms
- Allows you to set various range of functions(?) as loss

# XGboost

## XGBoost

$$\hat{y}_i^{(0)} = 0$$

$$\hat{y}_i^{(1)} = f_1(x_i) = \hat{y}_i^{(0)} + f_1(x_i)$$

$$\hat{y}_i^{(2)} = f_1(x_i) + f_2(x_i) = \hat{y}_i^{(1)} + f_2(x_i)$$

...

$$\hat{y}_i^{(t)} = \sum_{k=1}^t f_k(x_i) = \hat{y}_i^{(t-1)} + f_t(x_i)$$

**Composition structure**



$$\begin{aligned} \text{obj}^{(t)} &= \sum_{i=1}^n l(y_i, \hat{y}_i^{(t)}) + \sum_{i=1}^t \Omega(f_i) \\ &= \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t) + \text{constant} \end{aligned}$$

**Objective function**

$\Omega(f_i)$  — regularisation

# XGboost

## XGBoost

$$\begin{aligned}\text{obj}^{(t)} &= \sum_{i=1}^n l(y_i, \hat{y}_i^{(t)}) + \sum_{i=1}^t \Omega(f_i) \\ &= \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t) + \text{constant}\end{aligned}$$

**Objective function**

$$\begin{aligned}\text{obj}^{(t)} &= \sum_{i=1}^n (y_i - (\hat{y}_i^{(t-1)} + f_t(x_i)))^2 + \sum_{i=1}^t \Omega(f_i) \\ &= \sum_{i=1}^n [2(\hat{y}_i^{(t-1)} - y_i)f_t(x_i) + f_t(x_i)^2] + \Omega(f_t) + \text{constant}\end{aligned}$$

**MSE example**

$$\text{obj}^{(t)} = \sum_{i=1}^n [l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i)] + \Omega(f_t) + \text{constant}$$

$$g_i = \partial_{\hat{y}_i^{(t-1)}} l(y_i, \hat{y}_i^{(t-1)})$$

$$h_i = \partial_{\hat{y}_i^{(t-1)}}^2 l(y_i, \hat{y}_i^{(t-1)})$$

**2nd order Taylor's approximation**

# XGboost

## Regularisation

$$f_t(x) = w_{q(x)}, w \in R^T, q : R^d \rightarrow \{1, 2, \dots, T\}.$$

**Formal tree**

- $w$  — vector of scores on  $j$ -th index
- $q$  — function to assign object to leaf
- $T$  — number of leaves

$$\Omega(f) = \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2$$

**Complexity**

# XGboost

## Regularisation in learning

$$\begin{aligned}\text{obj}^{(t)} &\approx \sum_{i=1}^n [g_i w_{q(x_i)} + \frac{1}{2} h_i w_{q(x_i)}^2] + \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 \\ &= \sum_{j=1}^T [(\sum_{i \in I_j} g_i) w_j + \frac{1}{2} (\sum_{i \in I_j} h_i + \lambda) w_j^2] + \gamma T\end{aligned}$$

Obj function with Taylor and formal tree



$$\text{obj}^{(t)} = \sum_{j=1}^T [G_j w_j + \frac{1}{2} (H_j + \lambda) w_j^2] + \gamma T$$

$$G_j = \sum_{i \in I_j} g_i \quad H_j = \sum_{i \in I_j} h_i$$

«Simplification»

# XGboost

...

$$\begin{aligned}\text{obj}^{(t)} &\approx \sum_{i=1}^n [g_i w_{q(x_i)} + \frac{1}{2} h_i w_{q(x_i)}^2] + \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T w_j^2 \\ &= \sum_{j=1}^T [(\sum_{i \in I_j} g_i) w_j + \frac{1}{2} (\sum_{i \in I_j} h_i + \lambda) w_j^2] + \gamma T\end{aligned}$$

Obj function with Taylor and formal tree



$$\text{obj}^{(t)} = \sum_{j=1}^T [G_j w_j + \frac{1}{2} (H_j + \lambda) w_j^2] + \gamma T$$

$$G_j = \sum_{i \in I_j} g_i \qquad H_j = \sum_{i \in I_j} h_i$$

«Simplification»








# XGboost:

## Sorting schema:

$$Gain = \frac{1}{2} \left[ \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma$$

«Pruning»

 	  
g1, h1    g4, h4	g2, h2    g5, h5    g3, h3
$G_L = g_1 + g_4$	$G_R = g_2 + g_3 + g_5$

# XGboost

## Summing up:

- Decomposes loss-function in Taylor row
- Allows arbitrary set of loss functions
- Zip regularisation inside learning process
- Scan all dataset ones to find best split

# XGboost

## And his friends:

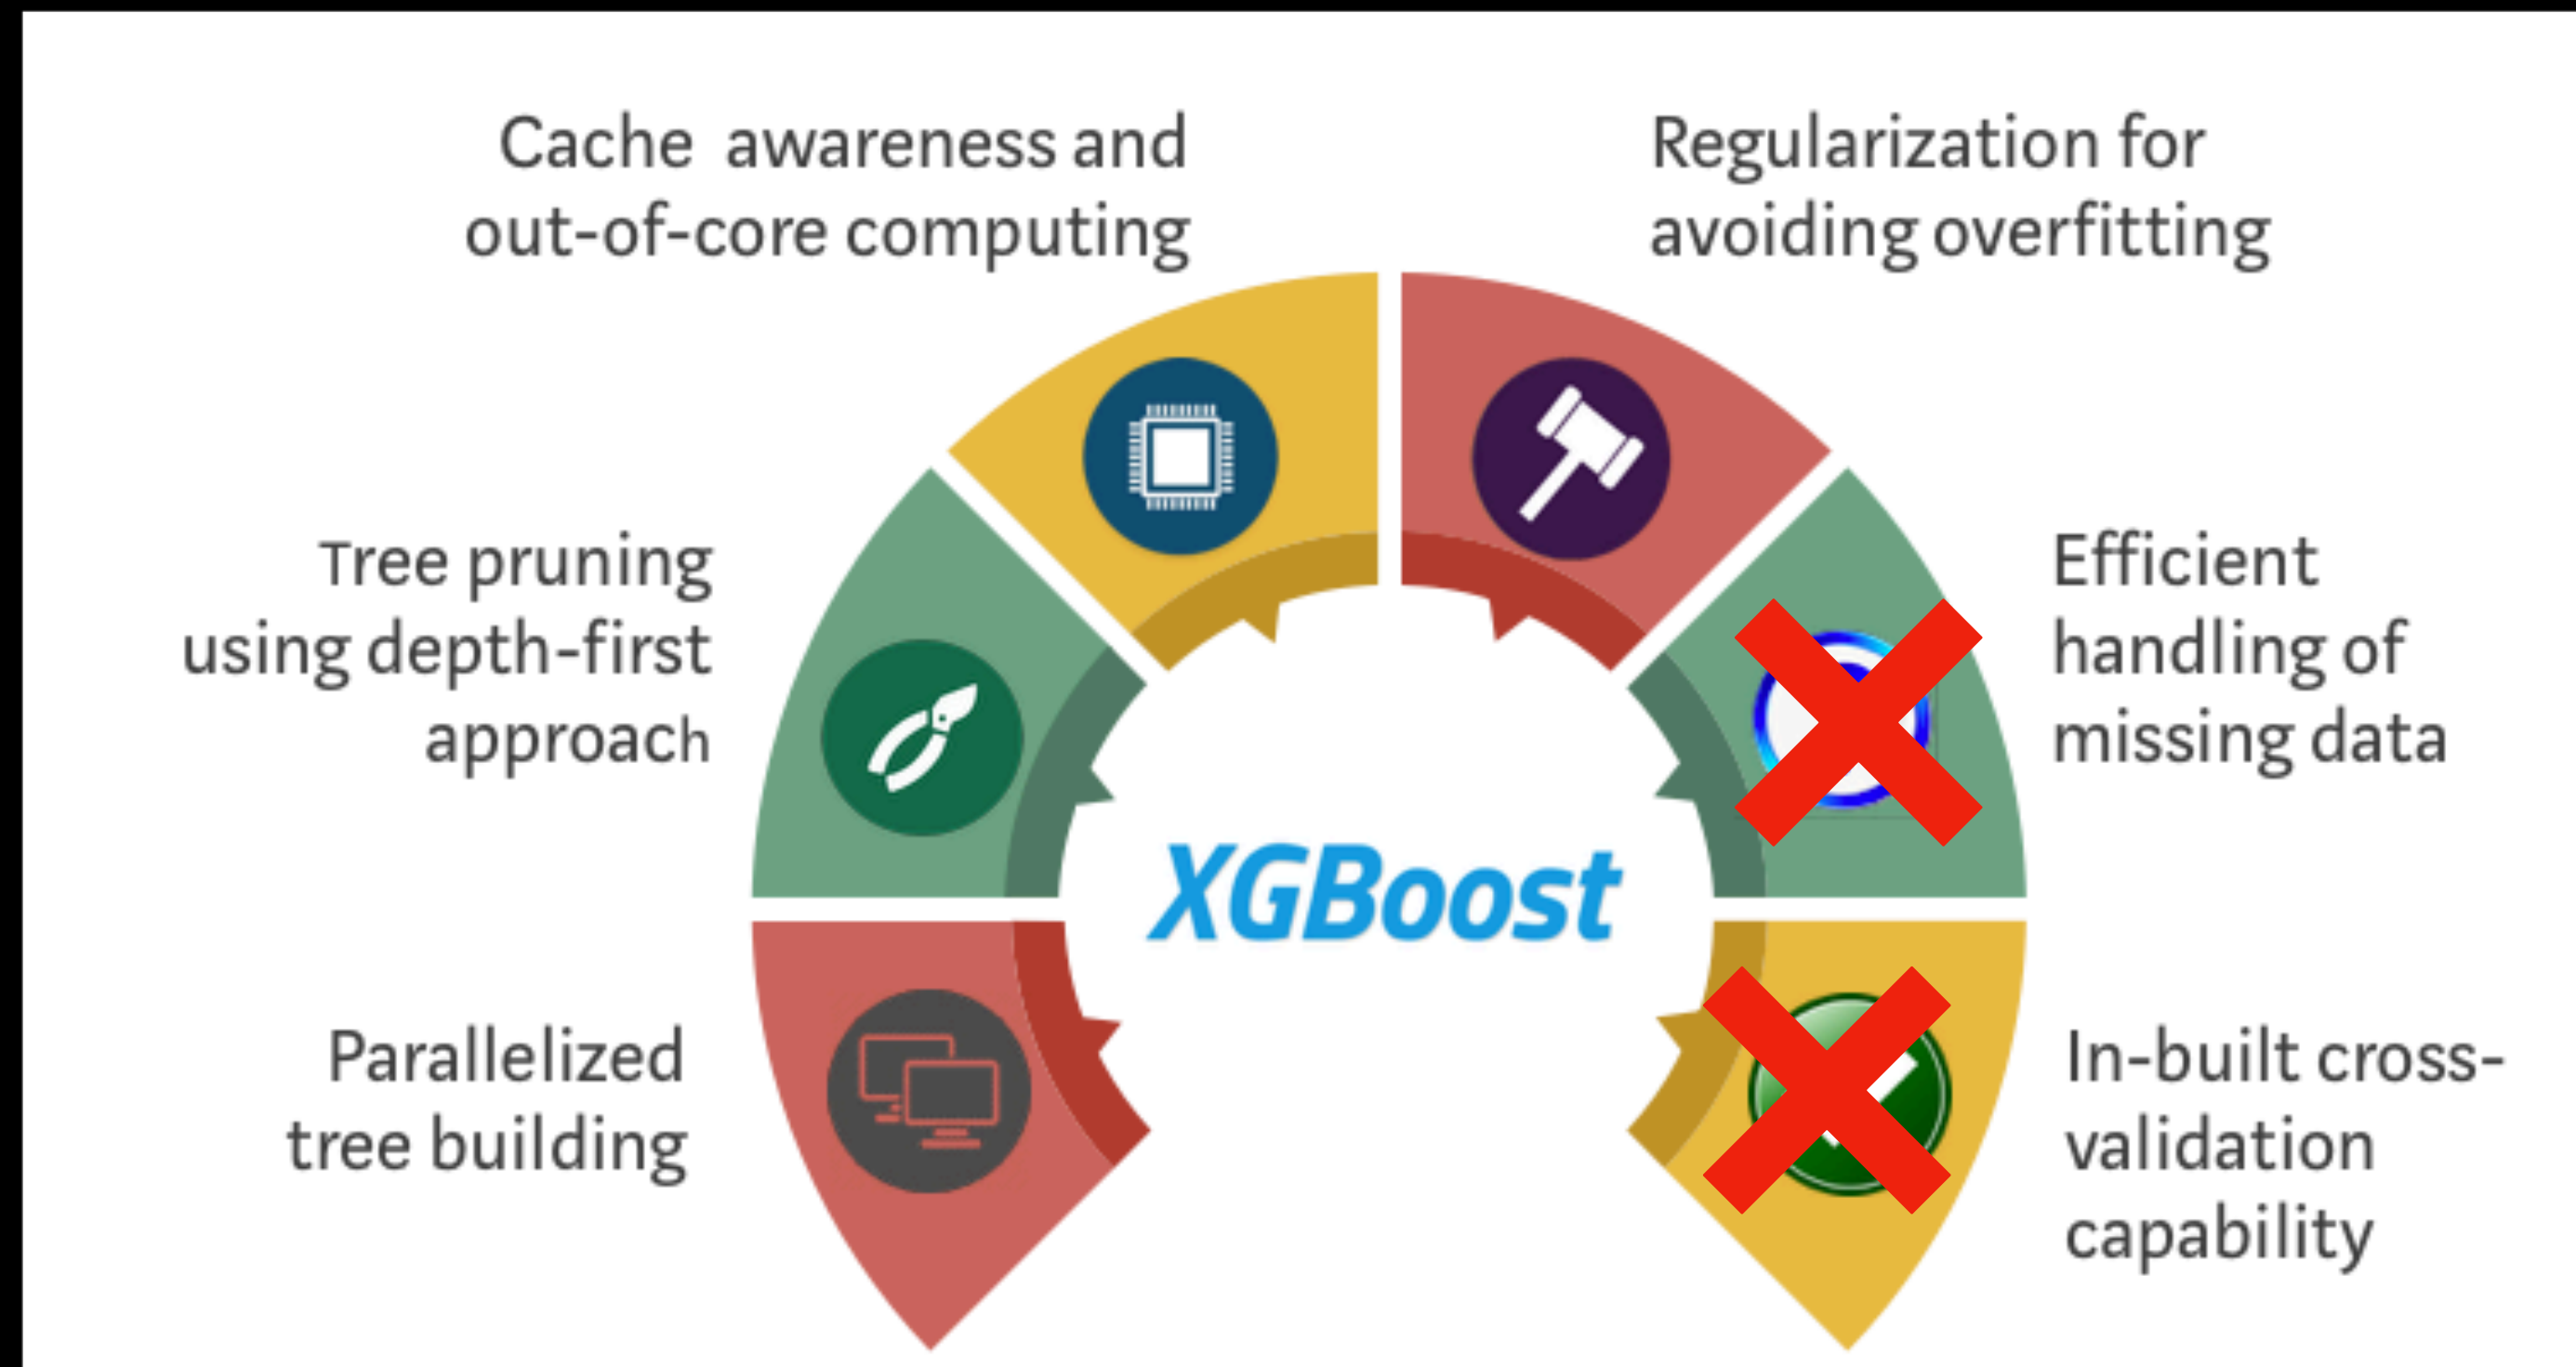
### Pros:

- Well-known (est. 2016)
- Available on ton different platforms (as cpp binding, mostly)
- Allows distribute training (spark, for example)

### Cons:

- Not so fancy comparing to IGBM, catBoost
- Not so powerful (~)
- Very strange sparsity

# XGBoost:



# XGBoost notes:

- XGBoost internal feature significance sucks, **use SHAP values**.
- XGBoost objectives:
  - **binary:logistic**: logistic regression for binary classification, output probability
  - **reg:squarederror**: regression with squared loss. Practically better on classification oO
  - **rank:{pairwise/map/ndcg}**: lambda mart ranking. Sometimes works significantly better than regr losses.
- Hyperparameters matters, tune them wisely (look refs).
- High gamma -> feature selection -> low gamma and features removed

# Learning to rank

## «Lambda»-smth

$$w = w + \eta \frac{\sigma}{1 + \exp(\sigma \langle x_i - x_j, w \rangle)} (x_i - x_j)$$

Gradient step optimising logit function on linear algorithm

$$w = w + \eta \frac{\sigma}{1 + \exp(\sigma \langle x_i - x_j, w \rangle)} |\Delta NDCG_{ij}| (x_i - x_j)$$

Gradient step optimising NDCG function on linear algorithm

NDCG delta  
by changing  $x_i$  and  $x_j$



# Learning to rank

## Summary

- Learning to rank - specific task with specific metrics and tricks.
- There are: point- pair- and list-wise algorithms.
- Best pair-wise algorithms looks as pointwise at runtime.
- There are tons ranking algorithms. Most powerful and/or used — GBDT
- XGBoost — allows arbitrary set of loss functions, incorporates regularisation into training, use Taylor 2nd order to estimate loss function.
- It's possible to optimise ranking metrics directly using lambda- techniques